

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report
import matplotlib.pyplot as plt
import seaborn as sns
import joblib

# Load dataset
df = pd.read_csv("/content/adult 3 (1).csv") # Rename your file accordingly

# Drop rows with missing values
df.replace('?', np.nan, inplace=True)
df.dropna(inplace=True)

# Drop unnecessary columns
df.drop(['education'], axis=1, inplace=True)

# Label encode categorical features
le = LabelEncoder()
for col in df.select_dtypes(include='object').columns:
    df[col] = le.fit_transform(df[col])

X = df.drop("income", axis=1)
y = df["income"]

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Scale features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Random Forest": RandomForestClassifier(),
    "KNN": KNeighborsClassifier(),
    "SVM": SVC(),
    "Gradient Boosting": GradientBoostingClassifier()
}

results = {}

for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    results[name] = acc
    print(f"{name}: {acc:.4f}")
    print(classification_report(y_test, y_pred))

```

Logistic Regression: 0.8194

	precision	recall	f1-score	support
0	0.84	0.94	0.89	2668
1	0.72	0.46	0.56	887
accuracy			0.82	3555
macro avg	0.78	0.70	0.72	3555
weighted avg	0.81	0.82	0.80	3555

Random Forest: 0.8459

	precision	recall	f1-score	support
0	0.88	0.93	0.90	2668
1	0.73	0.61	0.66	887
accuracy			0.85	3555
macro avg	0.80	0.77	0.78	3555
weighted avg	0.84	0.85	0.84	3555

```

KNN: 0.8233
      precision    recall  f1-score   support

     0       0.87       0.90       0.88       2668
     1       0.66       0.59       0.62        887

 accuracy         0.82         3555
 macro avg       0.77       0.75       0.75       3555
 weighted avg    0.82       0.82       0.82       3555

SVM: 0.8411
      precision    recall  f1-score   support

     0       0.86       0.94       0.90       2668
     1       0.75       0.54       0.63        887

 accuracy         0.84         3555
 macro avg       0.81       0.74       0.76       3555
 weighted avg    0.83       0.84       0.83       3555

Gradient Boosting: 0.8551
      precision    recall  f1-score   support

     0       0.88       0.94       0.91       2668
     1       0.77       0.60       0.68        887


 accuracy         0.86         3555
 macro avg       0.82       0.77       0.79       3555
 weighted avg    0.85       0.86       0.85       3555

```

```

plt.figure(figsize=(10, 5))
sns.barplot(x=list(results.keys()), y=list(results.values()))
plt.title("Model Accuracy Comparison")
plt.ylabel("Accuracy")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```



```

best_model_name = max(results, key=results.get)
best_model = models[best_model_name]
joblib.dump(best_model, "best_model.pkl")
print(f"Best model: {best_model_name} saved.")

```

```

Best model: Gradient Boosting saved.

```

```

from google.colab import files
files.download('best_model.pkl')

```

